

Daniel Esteban Cárdenas Jerez

Applied AI Engineer — Voice Agents, Multimodal Workflows & RAG

Brazil (Remote LATAM / Americas) · +57 318 884 8757

daniestebanc@hotmail.com · [linkedin.com/in/dancarjes](https://www.linkedin.com/in/dancarjes) · github.com/carjes232 · dev.danielcardenas.dev

Profile

Applied AI engineer with 3+ years of professional AI/ML experience and 5+ years using Python. Builds workflows that shorten process documentation, automate employee follow-up, and turn AI output into reviewable decisions and artifacts.

Experience

Lean Tech — AI Engineer · Medellín · Feb 2026 – Present

- Lead an agent workflow that turns process shadowing, source files, and revisions into editable SOPs and flowcharts, reducing manual drafting and keeping review in the same workflow.
- Led a voice agent that extends employee follow-up beyond longer absences and gives operations reviewable reported reasons; a voicemail-detection change cut observed cost per call **69%** (\$0.1275 to \$0.0395), implying **\$27.2K potential annual savings** at constant volume.
- Advised on an existing CRM intelligence product’s integrations and implemented focused Portal identity and delivery fixes; contributed a tested Microsoft Teams media-limit fix to OpenClaw.

Solenium — ML Engineer (Edge AI & Computer Vision) · Medellín · Jan 2025 – Feb 2026

- Led real-time energy-disaggregation systems with **sub-2-second inference** at 10 Hz.
- Deployed FaceID with liveness detection, OCR, offline handling, telemetry, APIs, CI/CD, and monitoring across **6+ devices**.

XM S.A. E.S.P — Python Data Engineer · Medellín · Jul 2024 – Jan 2025

- Built an auditable Python/pandas reconciliation workbench for **100M+ records**, reducing processing from **4–6 hours to <30 minutes**.

Magnus — Machine Learning Engineer (LLM, RAG & Azure) · Remote · Contract · May 2023 – Sep 2024

- Shipped Azure-hosted LLM workflows supporting **300+ daily clinical-simulation queries** with **<3s latency**.
- Built RAG and agents with FAISS, pgvector, Pinecone, LangGraph, tool calling, memory, FastAPI, recall@k evaluation, and quality dashboards.

Selected Projects

- **Multilingual Multi-Hop RAG:** FastAPI, BGE-M3, pgvector/HNSW, LangGraph, and Docker; source citations, +18% Recall@5, and 4.9-second P95 latency.
- **Document and presentation agents:** validated DOCX, flowcharts, and editable PPTX with source tracking, visual checks, and reviewable delivery; the Slides agent is in local qualification.
- **AI-Native Multi-Agent Engineering:** cost-aware routing, worktree isolation, durable callbacks, independent acceptance, and human-gated releases.

Technical Skills

Applied AI: RAG, agents, LangGraph, LlamaIndex, Azure OpenAI, Gemini, Ollama, tool calling, structured outputs, guardrails, computer vision

Retrieval & Evaluation: pgvector, Pinecone, FAISS, HNSW, BGE-M3, recall@k, citations, embeddings, latency and quality monitoring

Engineering: Python 3.11+, FastAPI, Flask, SQLAlchemy, TypeScript, React/TanStack, PostgreSQL, SQLite, Supabase, MCP, Docker, CI/CD, Cloudflare Workers, Azure/AWS

AgentOps/LLMOps: Multi-agent orchestration, Claude Code, OpenAI Codex, model routing, Git worktrees, tmux, durable callbacks, evaluation gates, human release controls

Education & Certifications

B.Sc. Engineering Physics — Universidad Nacional de Colombia · 2025 · IBM Generative AI Engineering with LLMs · IBM AI Engineering · IBM Deep Learning · AWS Academy Cloud Foundations

Languages: English — C1 Advanced · Spanish — Native · Portuguese — A2